



FINAL PROJECT DESCRIPTION

Automation Lab (540:383)

Fall 2024

I. Task description

Each group will be assigned one of two datasets: Housing Data (Regression Task) or Heart Failure Data (Classification Task). The task involves performing data analysis and modeling following the outline provided in Section 3. Deliverables include a Python notebook with the complete analysis, visualizations, and model implementation, along with a 10-slide PowerPoint presentation summarizing key insights and model performance.

Deadline: Sunday 12/08/2024

II. Goals

By completing this project, you will develop the ability to build and implement essential machine learning models, focusing on the two most common types in supervised learning: Regression and Classification. Additionally, you will enhance your data analysis skills by working with real-world datasets and further refine your Python programming expertise, particularly in the context of Data Science.

III. Project guidelines

1. Problem Definition

- Define the objective clearly. Understand the business problem or the research question you're trying to answer. This ensures the analysis is aligned with the goal.
- Define the type of task you want to perform: Regression/ Classification/Clustering and why?

2. Exploratory Data Analysis (EDA)

- Investigate the dataset through summary statistics, data visualization, and identifying relationships, patterns, or trends.
- Feature Classification
- Detect missing values, outliers, and distributions.
- Formulate hypotheses for further analysis.

3. Data Preprocessing

- Data Cleaning: Handle missing values, duplicates, outliers, and erroneous data.
- Data Transformation: Normalize, scale, or encode categorical variables.
- Feature Engineering: Create new features or variables that may improve the model.

4. Feature Selection

- Reduce the dimensionality by selecting the most relevant features based on domain knowledge or statistical techniques (e.g., correlation analysis, feature importance).
- Choose data for modeling (Input and Output)

5. Modeling

- Data Splitting: Split data into training, validation, and testing sets
- Choose appropriate machine learning models or algorithms based on the problem (e.g., regression, classification, clustering).
- Train the model using the training data.
- Obtain and show models, interpret the results

6. Model Evaluation

- Assess model performance using appropriate metrics (e.g., accuracy, precision, recall, F1-score for classification; RMSE, MAE for regression).
- Compare models based on evaluation metrics and select the best-performing model.
- Cross-validation to ensure the model generalizes well.

7. Model Interpretation and Insights

- Interpret the model's results and make sure they make sense in the context of the problem.
- Explain key drivers or features influencing the model.

IV. Grading rubric

- 7-Step Python Notebook & 10-Slide Presentation: 50%
- Error-Free Execution: 10%
- Consistency and Cohesiveness: 15%
- Accurate Interpretation and Conclusion: 15%
- Valuable Insights Extracted from Data: 10%

Good Luck!


```
from google.colab import drive
drive.mount('/content/gdrive')

file_path = '/content/gdrive/MyDrive/heart_failure data.csv'

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).
```

```
import pandas as pd
data= pd.read_csv(file_path)
data.head(5)
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium
0	75.0	0	582	0	20		1 265000.00	1.9	130
1	55.0	0	7861	0	38		0 263358.03	1.1	136
2	65.0	0	146	0	20		0 162000.00	1.3	129
3	50.0	1	111	0	20		0 210000.00	1.9	137
4	65.0	1	160	1	20		0 327000.00	2.7	116

Step 1: Problem Definition

Define the objective clearly. Understand the business problem or the research question you're trying to answer. This ensures the analysis is aligned with the goal.

Our objective is to understand what is causing a Death Event. By understanding what causes a Death Event it may help put a focus and emphasis on that specific issue. Understanding what causes Death Events can ultimately help reduce the number of Death Events.

Define the type of task you want to perform: Regression/ Classification/Clustering and why?

We will perform a classification analysis to determine how the independent variables serum creatinine, platelets, ejection fraction, serum sodium, and creatinine phosphokinase, impact the dependent variable Death Event. We are doing this to determine what events determine the change of a death event.

✓ Step 2: Exploratory Data Analysis (EDA)

```
#summary statistics
data.info()
data.describe()
```

```

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              299 non-null    float64
 1   anaemia          299 non-null    int64  
 2   creatinine_phosphokinase 299 non-null    int64  
 3   diabetes          299 non-null    int64  
 4   ejection_fraction 299 non-null    int64  
 5   high_blood_pressure 299 non-null    int64  
 6   platelets         299 non-null    float64
 7   serum_creatinine 299 non-null    float64
 8   serum_sodium      299 non-null    int64  
 9   sex               299 non-null    int64  
 10  smoking           299 non-null    int64  
 11  time              299 non-null    int64  
 12  DEATH_EVENT       299 non-null    int64  
dtypes: float64(3), int64(10)
memory usage: 30.5 KB

```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatini
count	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000
mean	60.833893	0.431438	581.839465	0.418060	38.083612	0.351171	263358.029264	1.393
std	11.894809	0.496107	970.287881	0.494067	11.834841	0.478136	97804.236869	1.034
min	40.000000	0.000000	23.000000	0.000000	14.000000	0.000000	25100.000000	0.500
25%	51.000000	0.000000	116.500000	0.000000	30.000000	0.000000	212500.000000	0.900
50%	60.000000	0.000000	250.000000	0.000000	38.000000	0.000000	262000.000000	1.100
75%	70.000000	1.000000	582.000000	1.000000	45.000000	1.000000	303500.000000	1.400
max	95.000000	1.000000	7861.000000	1.000000	80.000000	1.000000	850000.000000	9.400

data.shape

```

→ (299, 13)

```

```

# Import necessary libraries
data = pd.read_csv(file_path)

# Display basic information about the dataset
print("Dataset Info:")
print(data.info())
print("\nSummary Statistics:")
print(data.describe())

```

import matplotlib.pyplot as plt

```

→ Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              299 non-null    float64
 1   anaemia          299 non-null    int64  
 2   creatinine_phosphokinase 299 non-null    int64  
 3   diabetes          299 non-null    int64  
 4   ejection_fraction 299 non-null    int64  
 5   high_blood_pressure 299 non-null    int64  
 6   platelets         299 non-null    float64
 7   serum_creatinine 299 non-null    float64
 8   serum_sodium      299 non-null    int64  
 9   sex               299 non-null    int64  
 10  smoking           299 non-null    int64  
 11  time              299 non-null    int64  
 12  DEATH_EVENT       299 non-null    int64  
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
None

```

```

Summary Statistics:
      age    anaemia  creatinine_phosphokinase  diabetes \
count  299.000000  299.000000            299.000000  299.000000
mean   60.833893  0.431438            581.839465  0.418060

```

```

std    11.894809   0.496107          970.287881   0.494067
min    40.000000   0.000000          23.000000   0.000000
25%   51.000000   0.000000          116.500000   0.000000
50%   60.000000   0.000000          250.000000   0.000000
75%   70.000000   1.000000          582.000000   1.000000
max    95.000000   1.000000          7861.000000   1.000000

ejection_fraction  high_blood_pressure  platelets \
count            299.000000   299.000000   299.000000
mean             38.083612   0.351171   263358.029264
std              11.834841   0.478136   97804.236869
min              14.000000   0.000000   25100.000000
25%             30.000000   0.000000   212500.000000
50%             38.000000   0.000000   262000.000000
75%             45.000000   1.000000   303500.000000
max              80.000000   1.000000   850000.000000

serum_creatinine  serum_sodium   sex   smoking   time \
count            299.000000   299.000000   299.000000   299.000000
mean             1.39388   136.625418   0.648829   0.32107   130.260870
std              1.03451   4.412477   0.478136   0.46767   77.614208
min              0.50000   113.000000   0.000000   0.00000   4.000000
25%             0.90000   134.000000   0.000000   0.00000   73.000000
50%             1.10000   137.000000   1.000000   0.00000   115.000000
75%             1.40000   140.000000   1.000000   1.00000   203.000000
max              9.40000   148.000000   1.000000   1.00000   285.000000

DEATH_EVENT
count            299.00000
mean             0.32107
std              0.46767

```

Double-click (or enter) to edit

```

# Detect Missing Values
missing_values = data.isnull().sum()
print("\nMissing Values in Each Column:")
print(missing_values)

```

Missing Values in Each Column:

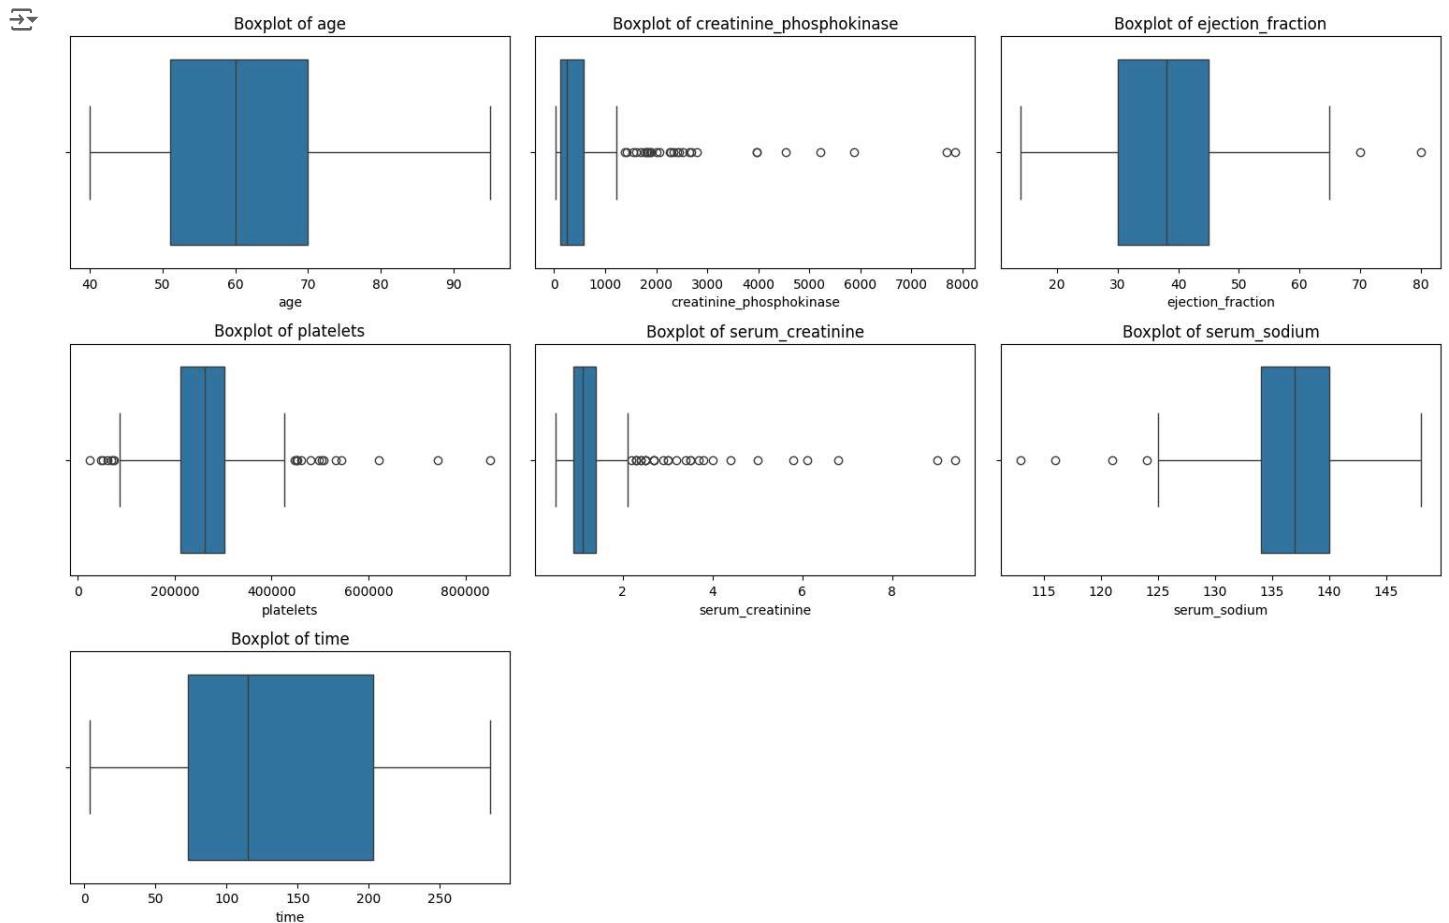
age	0
anaemia	0
creatinine_phosphokinase	0
diabetes	0
ejection_fraction	0
high_blood_pressure	0
platelets	0
serum_creatinine	0
serum_sodium	0
sex	0
smoking	0
time	0
DEATH_EVENT	0

dtype: int64

```

# Detect Outliers Using Boxplots for Numerical Features
numerical_features = ['age', 'creatinine_phosphokinase', 'ejection_fraction',
                      'platelets', 'serum_creatinine', 'serum_sodium', 'time']
import matplotlib.pyplot as plt
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(x=data[feature])
    plt.title(f"Boxplot of {feature}")
plt.xlabel(feature)
plt.tight_layout()
plt.show()

```

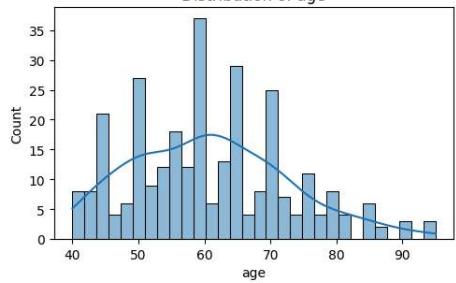


```
# Visualize Distributions of Numerical Features
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features, 1):
    plt.subplot(3, 3, i)
    sns.histplot(data[feature], bins=30, kde=True)
    plt.title(f"Distribution of {feature}")
    plt.xlabel(feature)
plt.tight_layout()
plt.show()

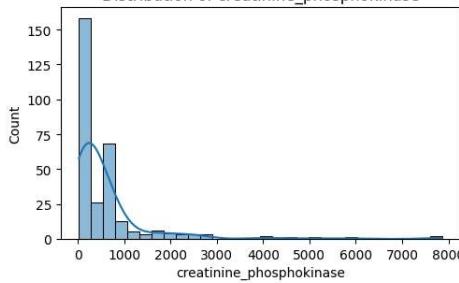
# Visualize Counts of Categorical Features
categorical_features = ['anaemia', 'diabetes', 'high_blood_pressure', 'sex', 'smoking', 'DEATH_EVENT']
plt.figure(figsize=(15, 10))
for i, feature in enumerate(categorical_features[:-1], 1): # Exclude DEATH_EVENT for now
    plt.subplot(3, 2, i)
    sns.countplot(x=data[feature])
    plt.title(f"Countplot of {feature}")
    plt.xlabel(feature)
plt.tight_layout()
plt.show()
```

[]

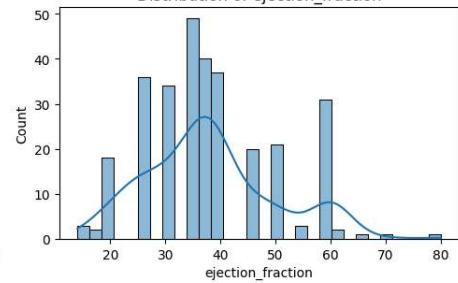
Distribution of age



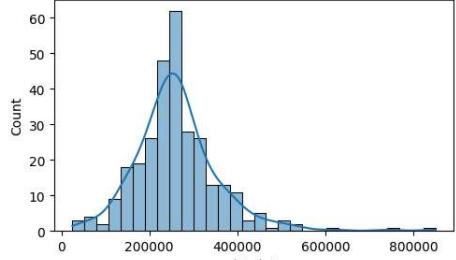
Distribution of creatinine_phosphokinase



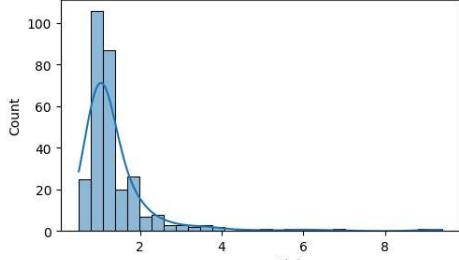
Distribution of ejection_fraction



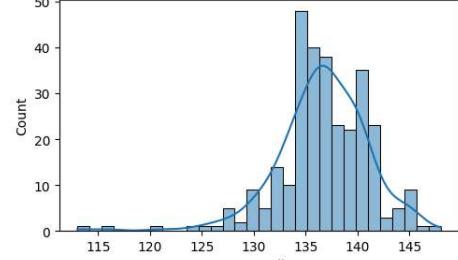
Distribution of platelets



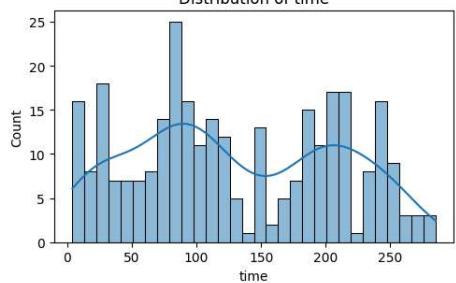
Distribution of serum_creatinine



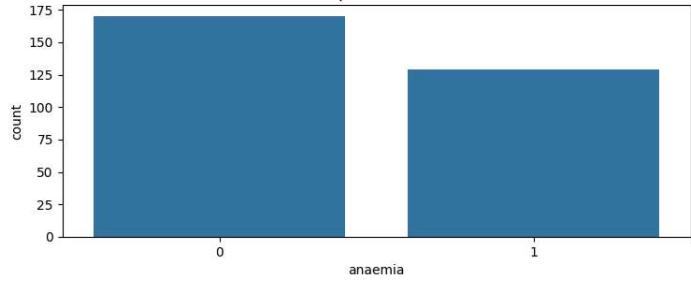
Distribution of serum_sodium



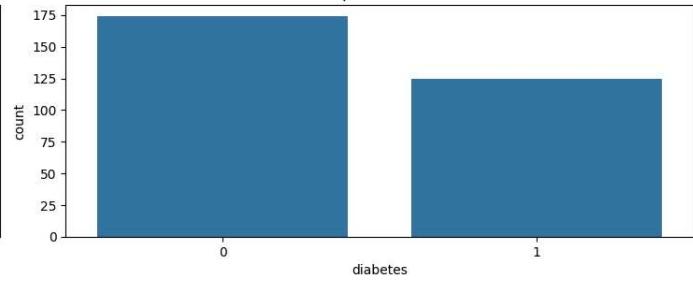
Distribution of time



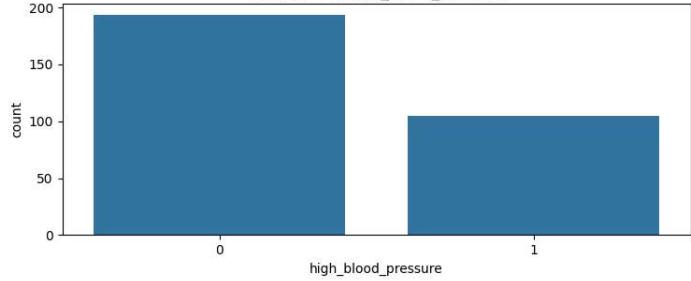
Countplot of anaemia



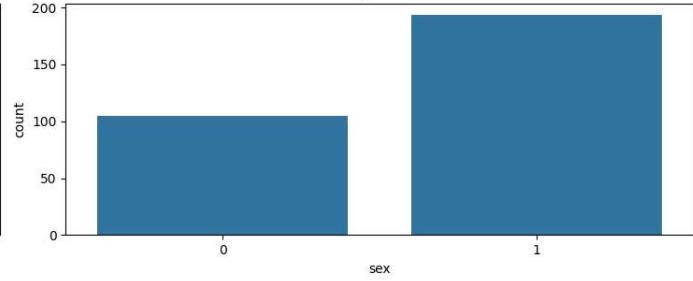
Countplot of diabetes



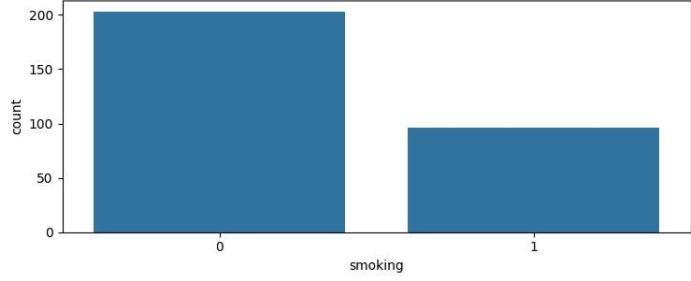
Countplot of high_blood_pressure



Countplot of sex



Countplot of smoking



feature selection:

```
#hypotheses for futher analysis
```

Null Hypothesis (H_0): There is no significant relationship between the categories (features such as age, anaemia, creatinine phosphokinase, diabetes, etc.) and the likelihood of a Death Event.

Alternative Hypothesis (H_1): At least one of the categories (features) significantly impacts the likelihood of a Death Event.

Feature Classification:

age: Numerical Continuous

anaemia: Categorical Nominal

creatinine_phosphokinase: Numerical Continuous

diabetes: Categorical Nominal

ejection_fraction: Numerical Discrete

high_blood_pressure: Categorical Nominal

platelets: Numerical Continuous

serum_creatinine: Numerical Continuous

serum_sodium: Numerical Continuous

sex: Categorical Nominal

smoking: Categorical Nominal

time: Numerical Time Series

DEATH_EVENT: Categorical Nominal

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
import matplotlib.pyplot as plt

# Fit decision tree classifier
clf = DecisionTreeClassifier(max_depth=3)
clf = clf.fit(X_train, y_train)

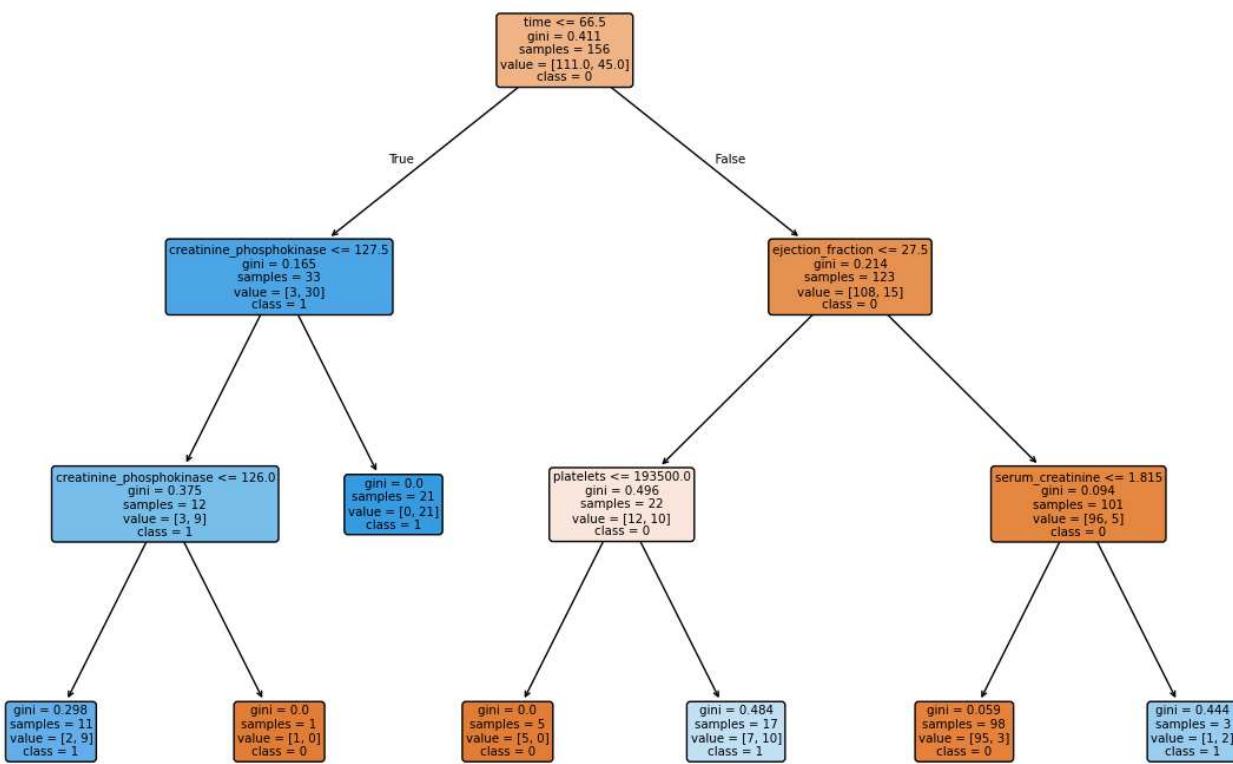
feature_names = X_train.columns
# Convert class names to strings
class_names = [str(cls) for cls in clf.classes_]

# Plot the decision tree
plt.figure(figsize=(15, 10))
tree.plot_tree(clf,
               filled=True,
               feature_names=feature_names,
```

```

class_names=class_names, # Use the converted class names
rounded=True)
plt.show()

```



If $\text{time} \leq 67.5$, the samples are sent to the left branch. If $\text{time} > 67.5$, the samples are sent to the right branch.

Left branch 1: If $\text{ejection fraction} \leq 72.5$, the samples are further split into the next level. If $\text{ejection fraction} > 72.5$, only one patient is classified, with no death event (class 0).

Left branch 2: If $\text{creatinine phosphokinase} \leq 109.5$, the samples are split further. If $\text{creatinine phosphokinase} > 109.5$, there is a strong majority of no death events (class 0)

Right branch 1: If $\text{serum creatinine} \leq 1.815$, the majority of patients experience no death events (class 0). If $\text{serum creatinine} > 1.815$, there's a higher proportion of death events, leading to further splits.

Right branch 2: If $\text{platelets} \leq 90,000$, the data splits again based on serum sodium . If $\text{platelets} > 90,000$, the model predicts no death events (class 0) with very high confidence (low Gini index of 0.093).

Right branch 3: If $\text{serum sodium} \leq 132.5$, the risk of a death event (class 1) is high. If $\text{serum sodium} > 132.5$, the risk of a death event is lower.

✓ Step 3: Data Preprocessing

```

# Load the dataset
data = pd.read_csv(file_path)

# 1. Detect Missing Values
print("Missing Values in Each Column:")
missing_values = data.isnull().sum()
print(missing_values)
  
```

```

→ Missing Values in Each Column:
age 0
anaemia 0
creatinine_phosphokinase 0
diabetes 0
ejection_fraction 0
high_blood_pressure 0
platelets 0
serum_creatinine 0
serum_sodium 0
sex 0
smoking 0
time 0
DEATH_EVENT 0
dtype: int64

# 2. Handle Duplicates
duplicates = data.duplicated().sum()
print(f"\nNumber of duplicate rows: {duplicates}")

# If duplicates exist, remove them
data = data.drop_duplicates()

→
Number of duplicate rows: 0

# 3. Detect and Handle Outliers
# Use IQR to filter out outliers in numerical features
numerical_features = ['age', 'creatinine_phosphokinase', 'ejection_fraction',
                      'platelets', 'serum_creatinine', 'serum_sodium', 'time']

for feature in numerical_features:
    Q1 = data[feature].quantile(0.25)
    Q3 = data[feature].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Filter out outliers
    outliers_removed = data[(data[feature] >= lower_bound) & (data[feature] <= upper_bound)]

    print(f"\nFeature: {feature}")
    print(f"Outliers removed: {len(data) - len(outliers_removed)}")

    # Update the dataset by removing outliers
    data = outliers_removed

→
Feature: age
Outliers removed: 0

Feature: creatinine_phosphokinase
Outliers removed: 29

Feature: ejection_fraction
Outliers removed: 2

Feature: platelets
Outliers removed: 18

Feature: serum_creatinine
Outliers removed: 23

Feature: serum_sodium
Outliers removed: 3

Feature: time
Outliers removed: 0

# 4. Check for Erroneous Data (e.g., negative values where not applicable)
for feature in numerical_features:
    if (data[feature] < 0).any():
        print(f"Erroneous data found in {feature}")

# Display cleaned dataset info
print("\nCleaned Dataset Info:")
print(data.info())

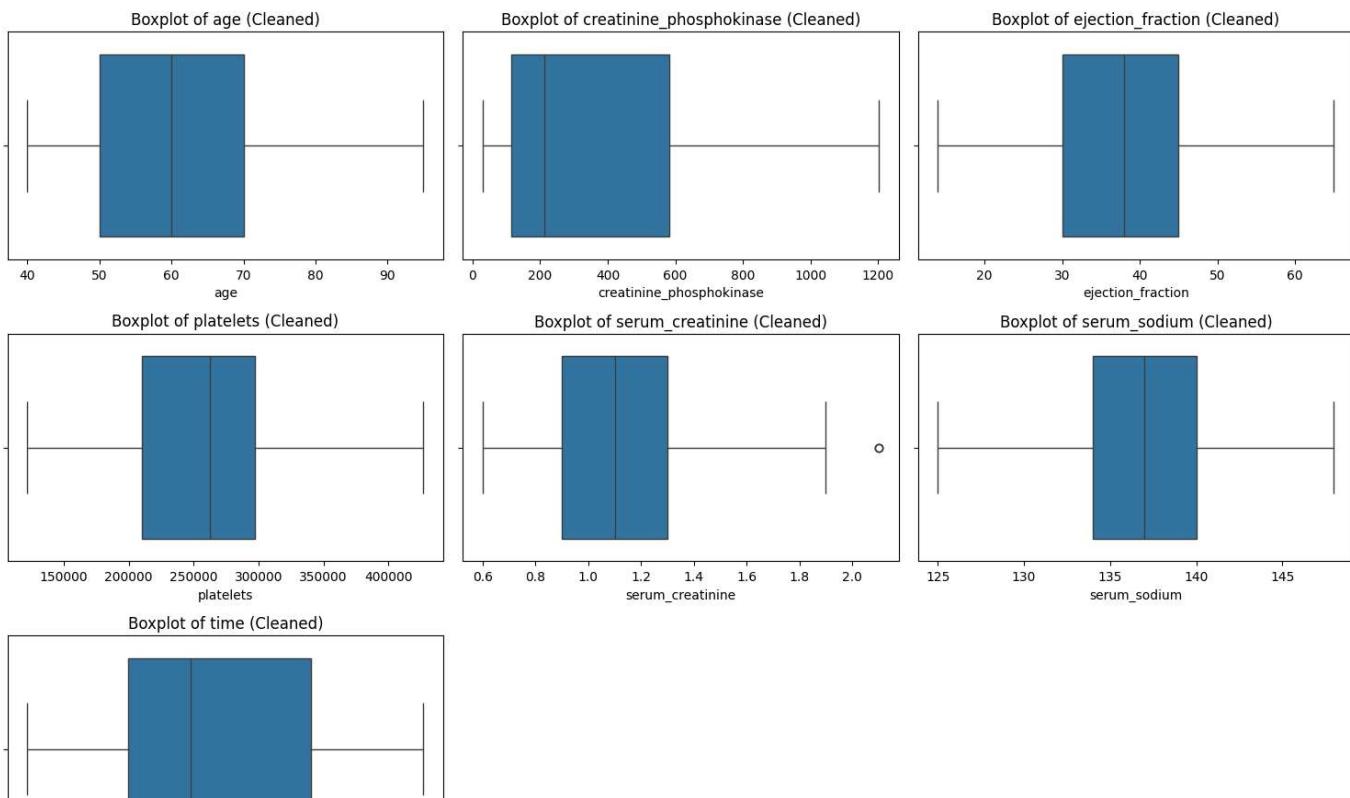
```

```
# Visualize to confirm cleaning process
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(x=data[feature])
    plt.title(f"Boxplot of {feature} (Cleaned)")
plt.tight_layout()
plt.show()
```



Cleaned Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
Index: 224 entries, 0 to 298
Data columns (total 13 columns):
 #   Column           Non-Null Count Dtype  
 ---  -- 
 0   age              224 non-null   float64 
 1   anaemia          224 non-null   int64   
 2   creatinine_phosphokinase 224 non-null   int64  
 3   diabetes          224 non-null   int64  
 4   ejection_fraction 224 non-null   int64  
 5   high_blood_pressure 224 non-null   int64  
 6   platelets         224 non-null   float64 
 7   serum_creatinine  224 non-null   float64 
 8   serum_sodium      224 non-null   int64  
 9   sex               224 non-null   int64  
 10  smoking           224 non-null   int64  
 11  time              224 non-null   int64  
 12  DEATH_EVENT       224 non-null   int64  
dtypes: float64(3), int64(10)
memory usage: 24.5 KB
None
```



```
# Feature Engineering: Create new features or variables that may improve the model
# Example: Create a feature for age group
data['age_group'] = pd.cut(data['age'], bins=[0, 50, 70, 100], labels=['young', 'middle_aged', 'old'])

# Encode the new feature
age_group_encoded = pd.get_dummies(data['age_group'], drop_first=True)

# Concatenate with the original dataset
data = pd.concat([data.drop(columns=['age_group']), age_group_encoded], axis=1)

print("Feature Engineered Data:")
print(data.head())
```



Feature Engineered Data:

```
age  anaemia  creatinine_phosphokinase  diabetes  ejection_fraction  \

```

```

0 75.0      0          582      0        20
2 65.0      0          146      0        20
3 50.0      1          111      0        20
5 90.0      1          47       0        40
6 75.0      1          246      0        15

  high_blood_pressure  platelets  serum_creatinine  serum_sodium  sex \
0                  1  265000.0           1.9         130      1
2                  0  162000.0           1.3         129      1
3                  0  210000.0           1.9         137      1
5                  1  204000.0           2.1         132      1
6                  0  127000.0           1.2         137      1

smoking  time  DEATH_EVENT  middle_aged  old
0       0     4           1      False    True
2       1     7           1      True   False
3       0     7           1      False   False
5       1     8           1      False    True
6       0    10          1      False    True

```

✓ Step 4: Feature Selection

```

from sklearn.model_selection import train_test_split

# Define features
X = data.drop(['DEATH_EVENT'], axis=1)
y = data['DEATH_EVENT']

```

✓ Step 5: Modeling

```

# Split the data: 70% training, 30% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Show the shape of the resulting datasets
print(f"X_train shape: {X_train.shape}, y_train shape: {y_train.shape}")
print(f"X_test shape: {X_test.shape}, y_test shape: {y_test.shape}")

→ X_train shape: (156, 14), y_train shape: (156,)
    X_test shape: (68, 14), y_test shape: (68,)

```

```

# Train a Decision Tree Classifier
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)

```

DecisionTreeClassifier(max_depth=3, random_state=42)

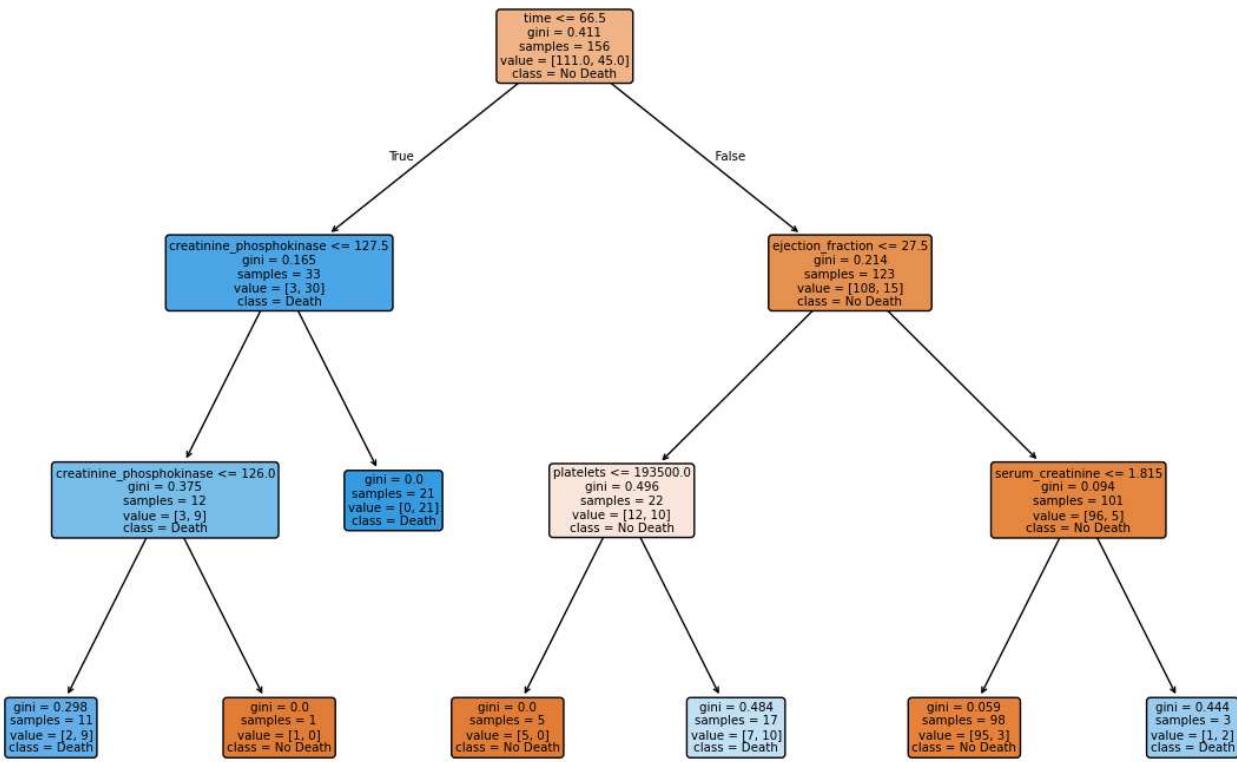
```

from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
# Visualize the Decision Tree
plt.figure(figsize=(15, 10))
plot_tree(clf, filled=True, feature_names=X.columns, class_names=['No Death', 'Death'], rounded=True)
plt.title("Decision Tree Visualization")
plt.show()

```



Decision Tree Visualization



```
from sklearn.metrics import classification_report, accuracy_score
# Predict on test data
y_pred = clf.predict(X_test)
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

→ Model Accuracy: 0.79

Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.85	0.86	52
1	0.56	0.62	0.59	16
accuracy			0.79	68
macro avg	0.72	0.74	0.73	68
weighted avg	0.80	0.79	0.80	68

✓ Step 6: Model Evaluation

Confusion Matrix

```

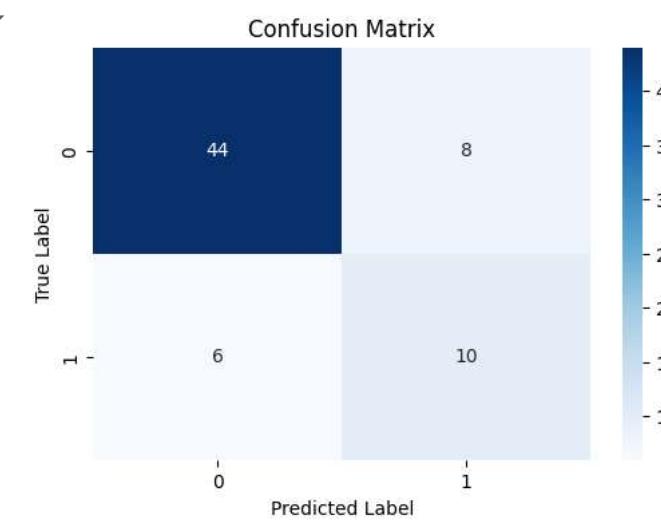
#Confusion Matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
y_pred = clf.predict(X_test)
# Compute the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Create a heatmap to display the binary confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

# Add labels and title
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')

# Show the plot
plt.show()

```



Accuracy Matrix

```

#Accuracy Metric
from sklearn.metrics import accuracy_score

# Calculate accuracy score
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

```

Model Accuracy: 0.79

ROC-AUC

```

#ROC-AUC
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
y_pred = clf.predict(X_test)
fpr, tpr, thresholds = roc_curve(y_test,y_pred)
roc_auc= auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color= 'darkorange', lw=2, label='ROC curve (area= %0.2f)' %roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```

<Automation Lab>

<Fall 2024>

Heart

Failure

Data Set

Olivia R., Somto O., Zach G.



<Step 1>

<Problem Definition>

Objective:

To understand what is causing a Death Event. By understanding what causes a Death Event it may help put a focus and emphasis on that specific issue. Understanding what causes Death Events can ultimately help reduce the number of Death Events.

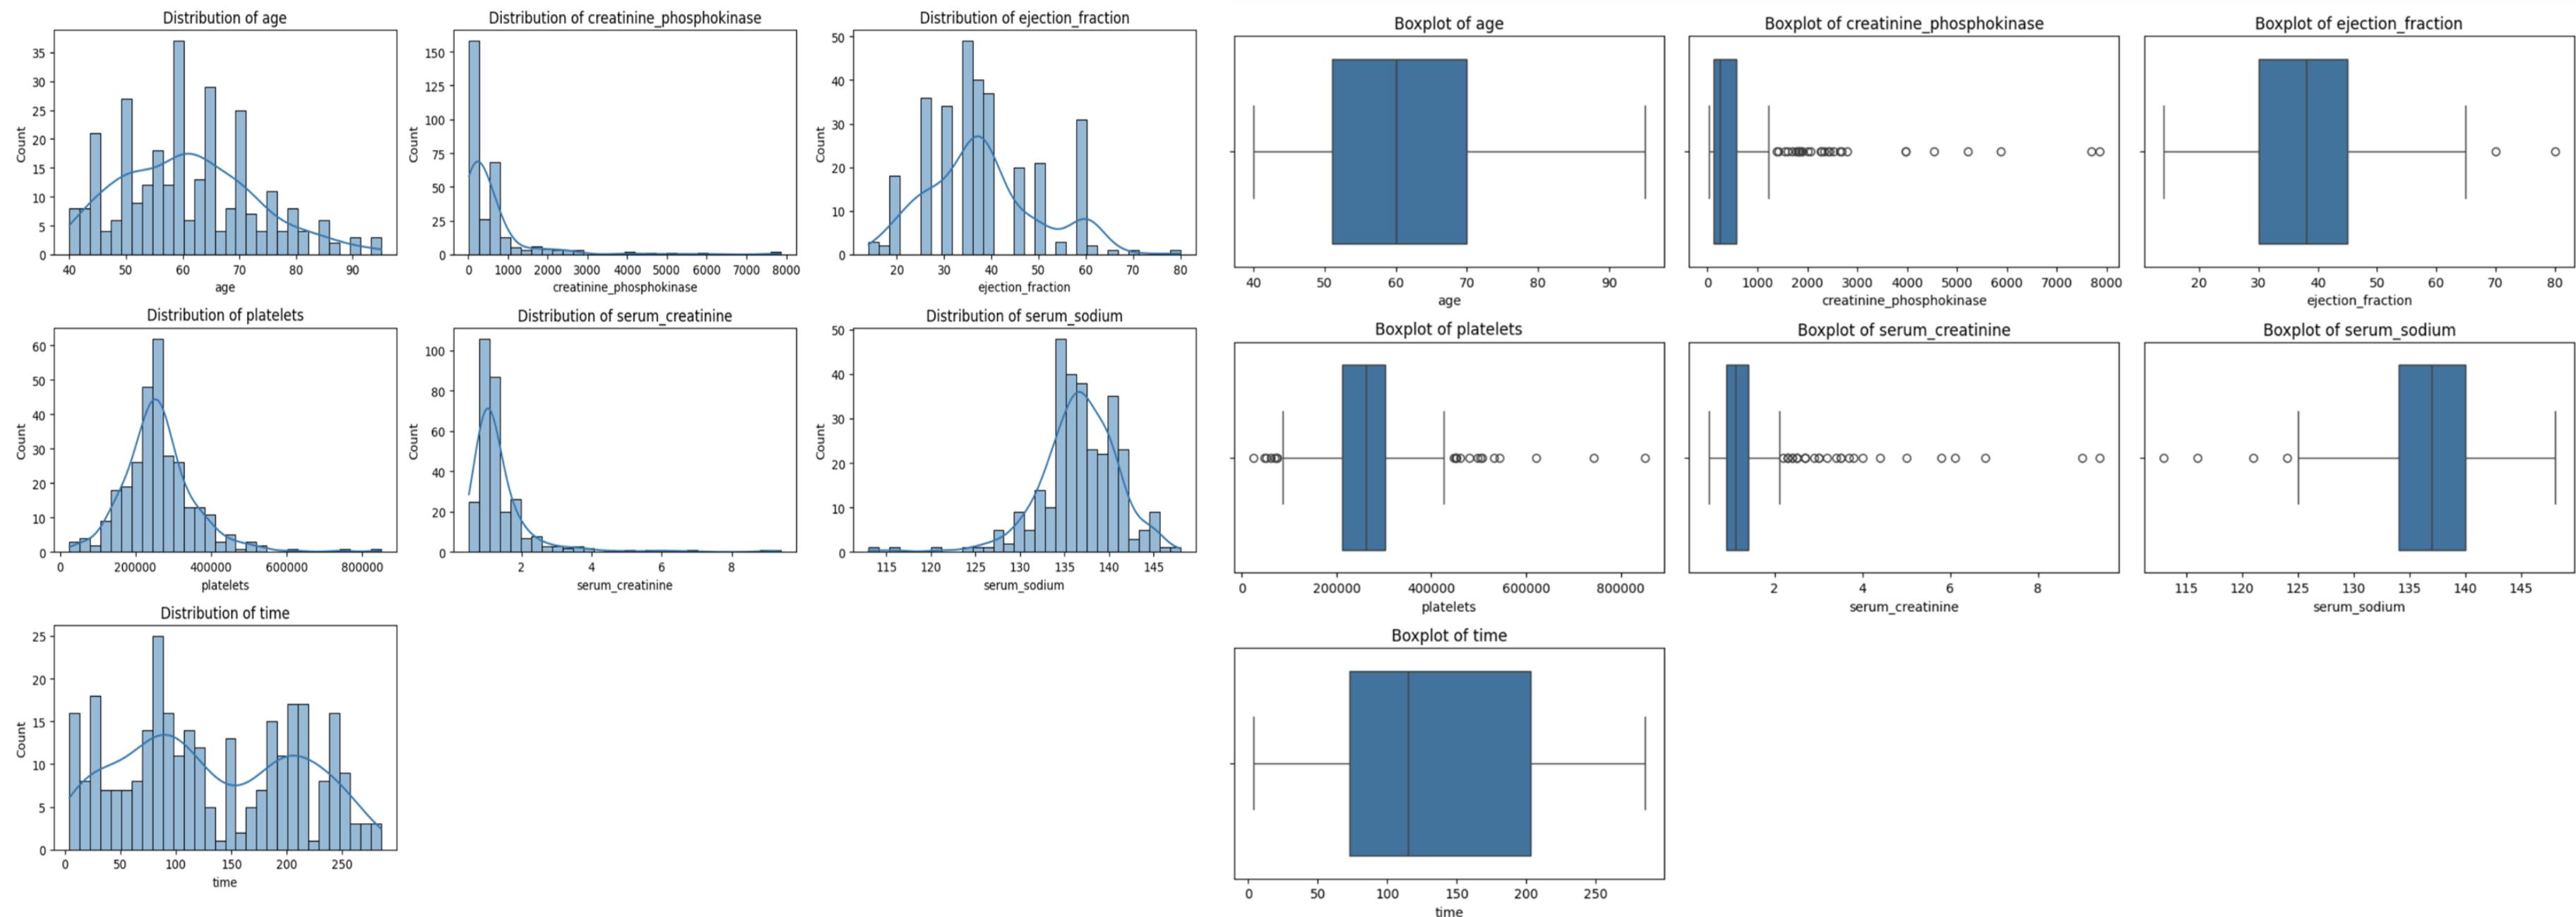
Task:

We will perform a classification analysis to determine how the independent variables impact the dependent variable Death Event. We are doing this to determine what events determine the change of a death event.

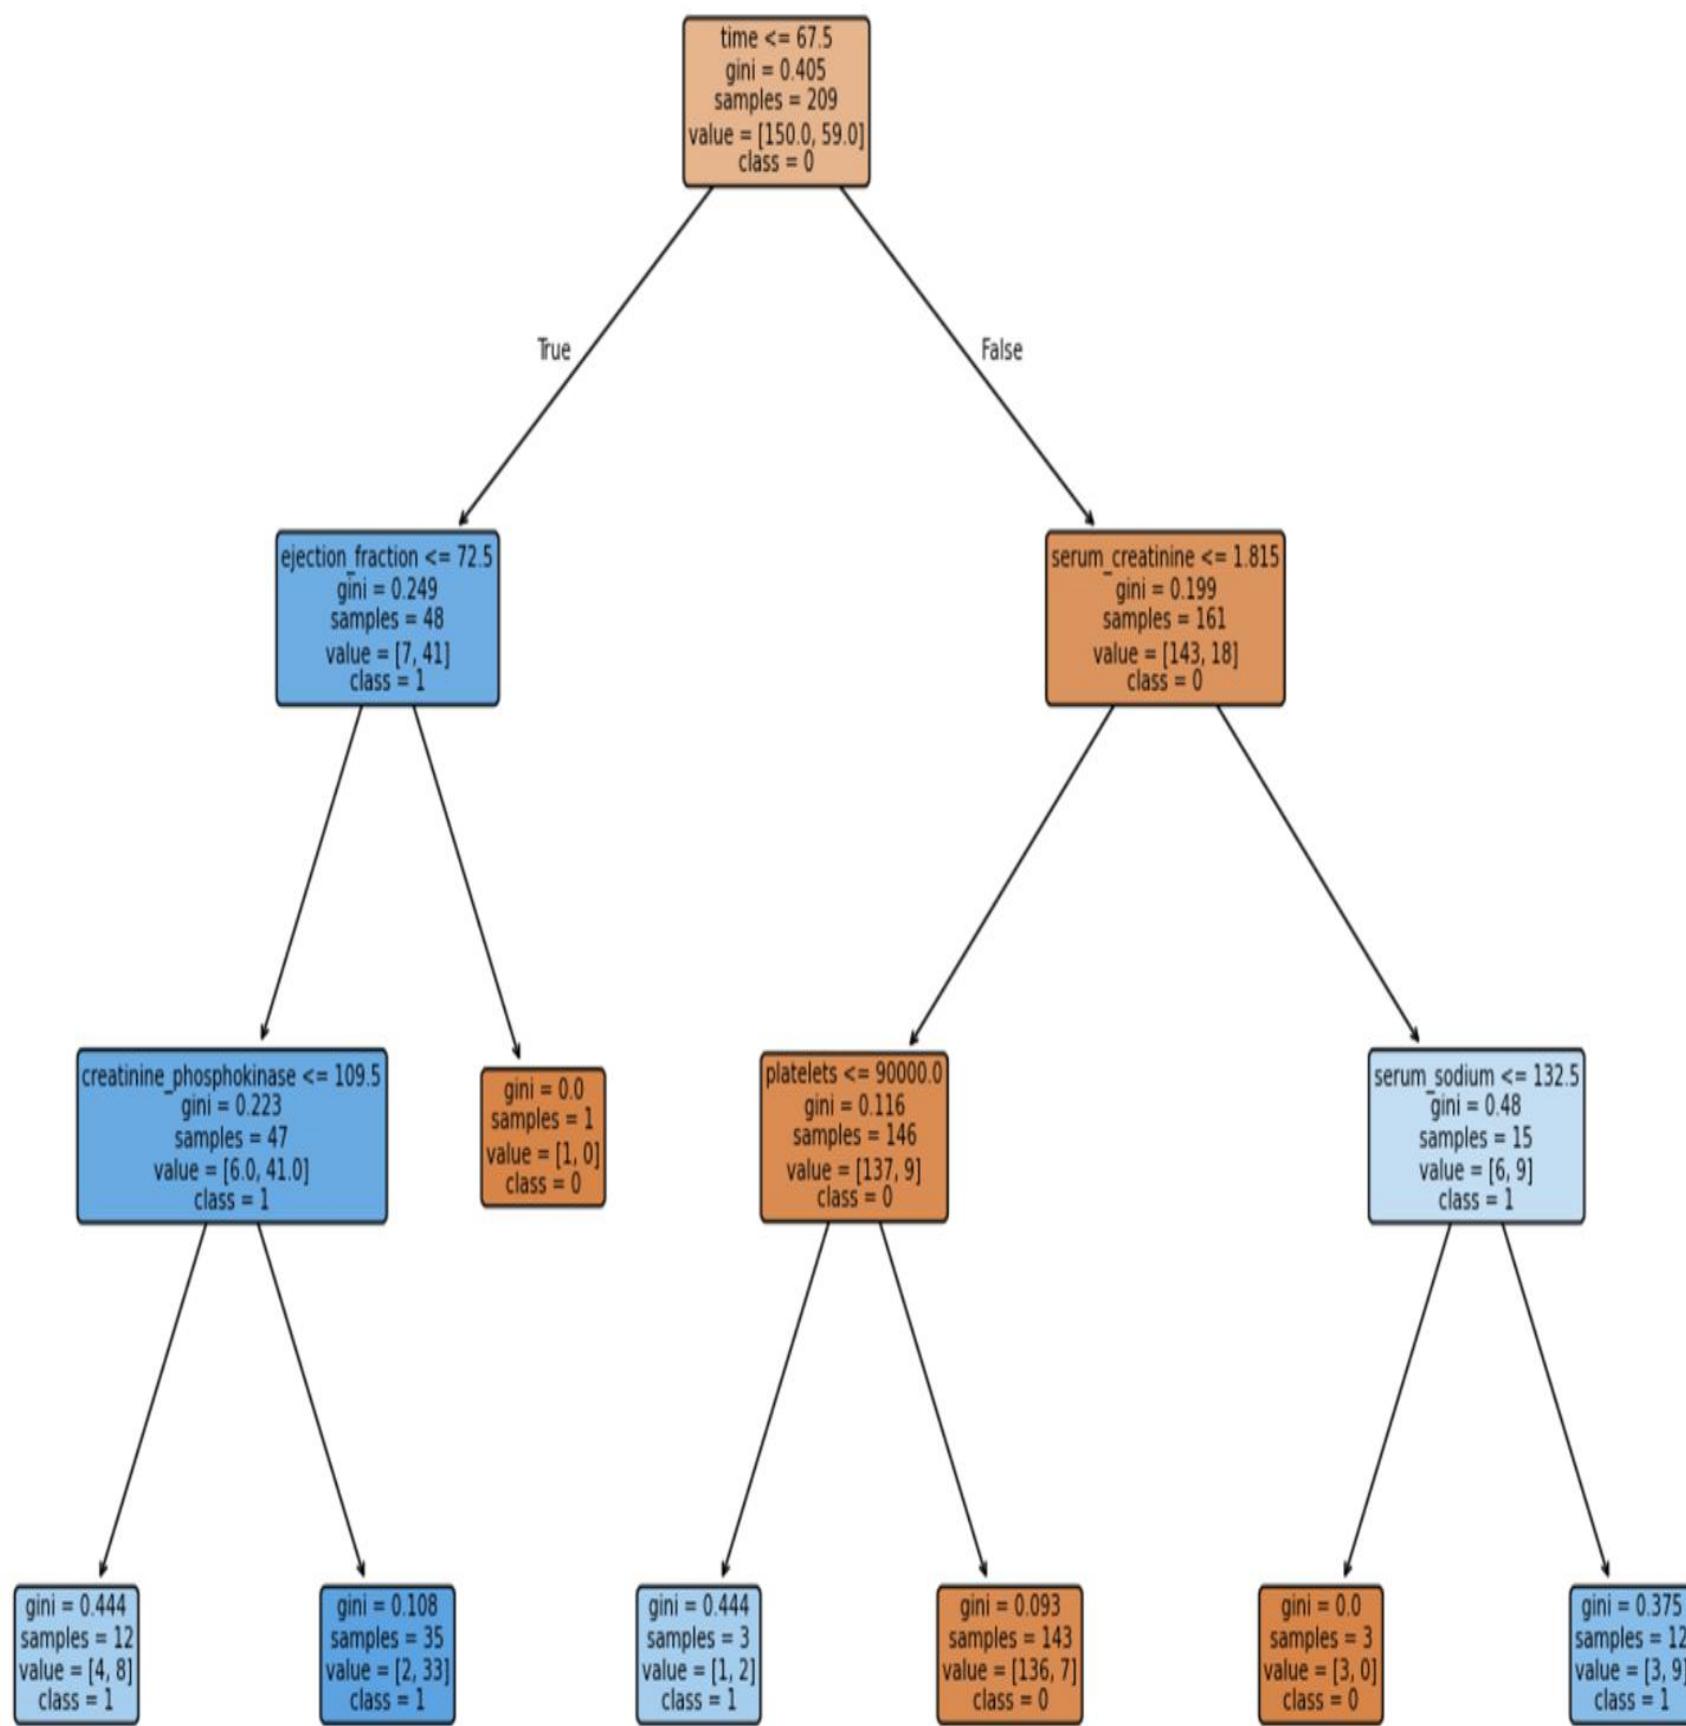
<Step 2>

<Exploratory Data Analysis (EDA)>

- Zero missing values detected
- Creatinine Phosphokinase and Serum Creatinine have the most outliers and have a right skewed graph
- Null Hypothesis (H_0): There is no significant relationship between the categories (features such as age, anaemia, creatinine phosphokinase, diabetes, etc.) and the likelihood of a Death Event.
- Alternative Hypothesis (H_1): At least one of the categories (features) significantly impacts the likelihood of a Death Event.



<Step 2>



<Exploratory Data Analysis (EDA)>

If $\text{time} \leq 67.5$, the samples are sent to the left branch.
If $\text{time} > 67.5$, the samples are sent to the right branch.

Left branch 1:

If $\text{ejection fraction} \leq 72.5$, the samples are further split into the next level.
If $\text{ejection fraction} > 72.5$, only one patient is classified, with no death event (class 0).

Left branch 2:

If $\text{serum creatinine} \leq 1.815$, the samples are split further.
If $\text{serum creatinine} > 1.815$, there is a strong majority of no death events (class 0)

Right branch 1:

If $\text{creatinine phosphokinase} \leq 109.5$, the majority of patients experience no death events (class 0).

If $\text{creatinine phosphokinase} > 109.5$, there's a higher proportion of death events, leading to further splits.

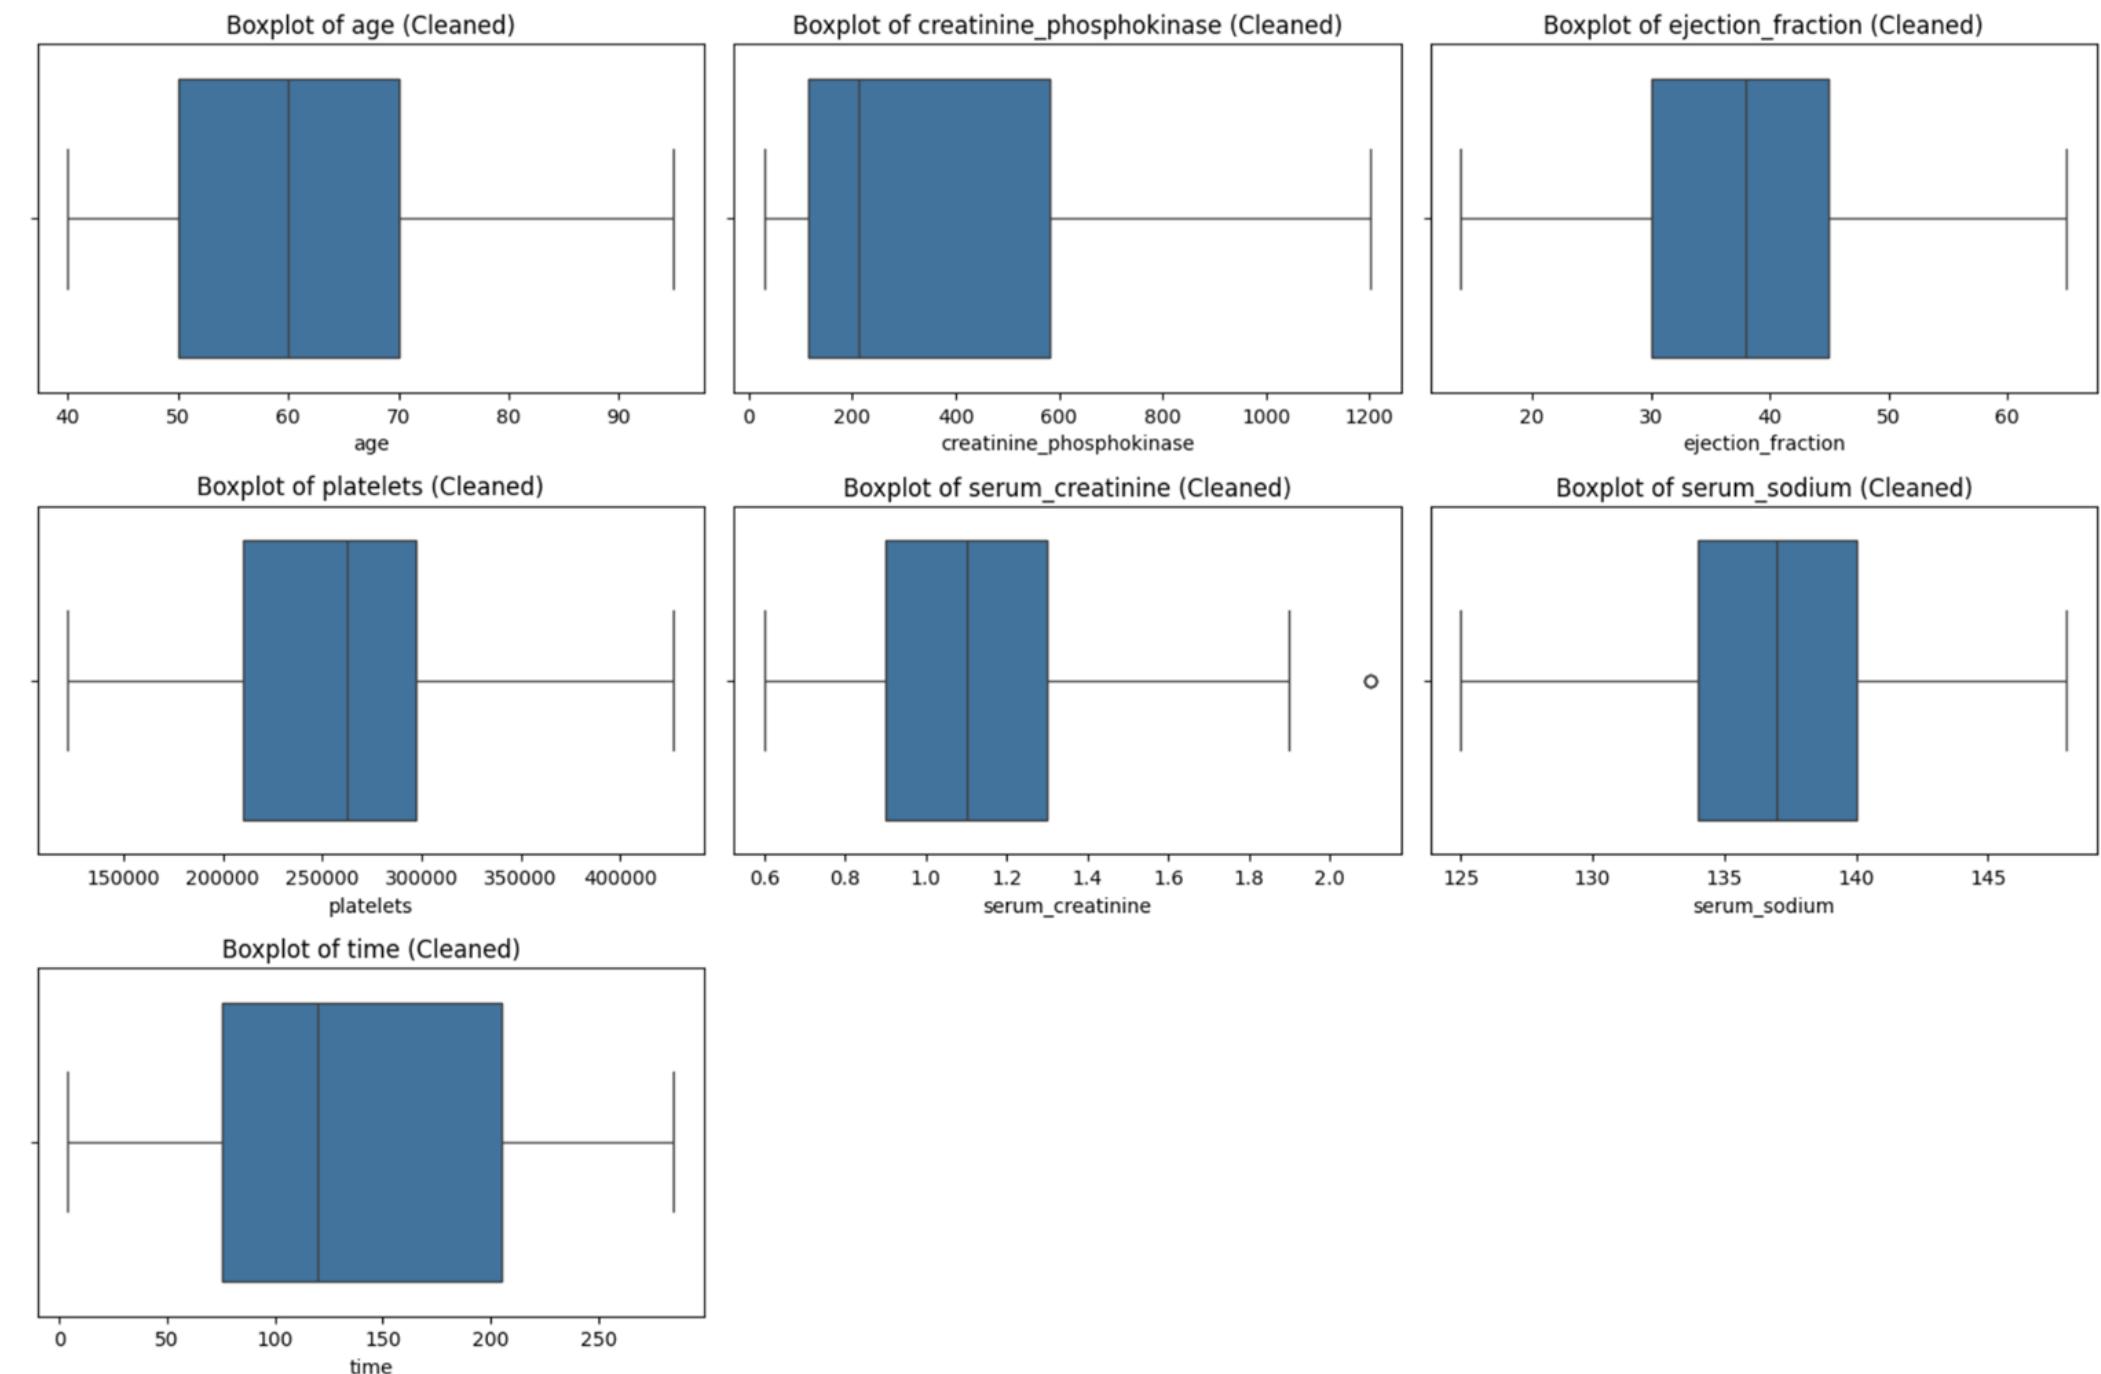
Right branch 2:

If $\text{serum sodium} \leq 132.5$, the risk of a death event (class 1) is high.
If $\text{serum sodium} > 132.5$, the risk of a death event is lower.

<Step 3>

<Data Preprocessing>

- Dropped any duplicates
- Detected and handled outliers by removing them front the dataset
 - 29 from creatinine_phosphokinase
 - 2 from ejection_fraction
 - 18 from platelets
 - 23 from serum_creatinine
 - 3 from serum_sodium
- Check for Erroneous Data (negative values where not applicable)
- Create a feature for age group



<Step 4>

<Feature Selection>

```
X = data.drop(['DEATH_EVENT'], axis=1)  
y = data['DEATH_EVENT']
```

```
X_train shape: (209, 12), y_train shape: (209, )  
X_test shape: (90, 12), y_test shape: (90, )
```

Split Data into Features and Target

- Features (**X**): All columns except DEATH_EVENT.
- Target (**y**): The DEATH_EVENT column.

Dropped the Target Variable (**DEATH_EVENT**) from Features to avoid data leakage.

Used the **train_test_split Function** from `sklearn.model_selection` to separate data into training and testing sets.

<Step 5>

<Modeling>

Data Splitting:

- Divided the dataset into **70% training** and **30% testing** using `train_test_split`.
- Maintained reproducibility with `random_state=42`.

Model Training:

- Used a **Decision Tree Classifier** with a maximum depth of 3 to avoid overfitting.
- Trained the model on the **training set (`X_train`, `y_train`)**.

Model Visualization:

- Plotted and visualized the **Decision Tree**, showing key splits and decision paths.

Prediction:

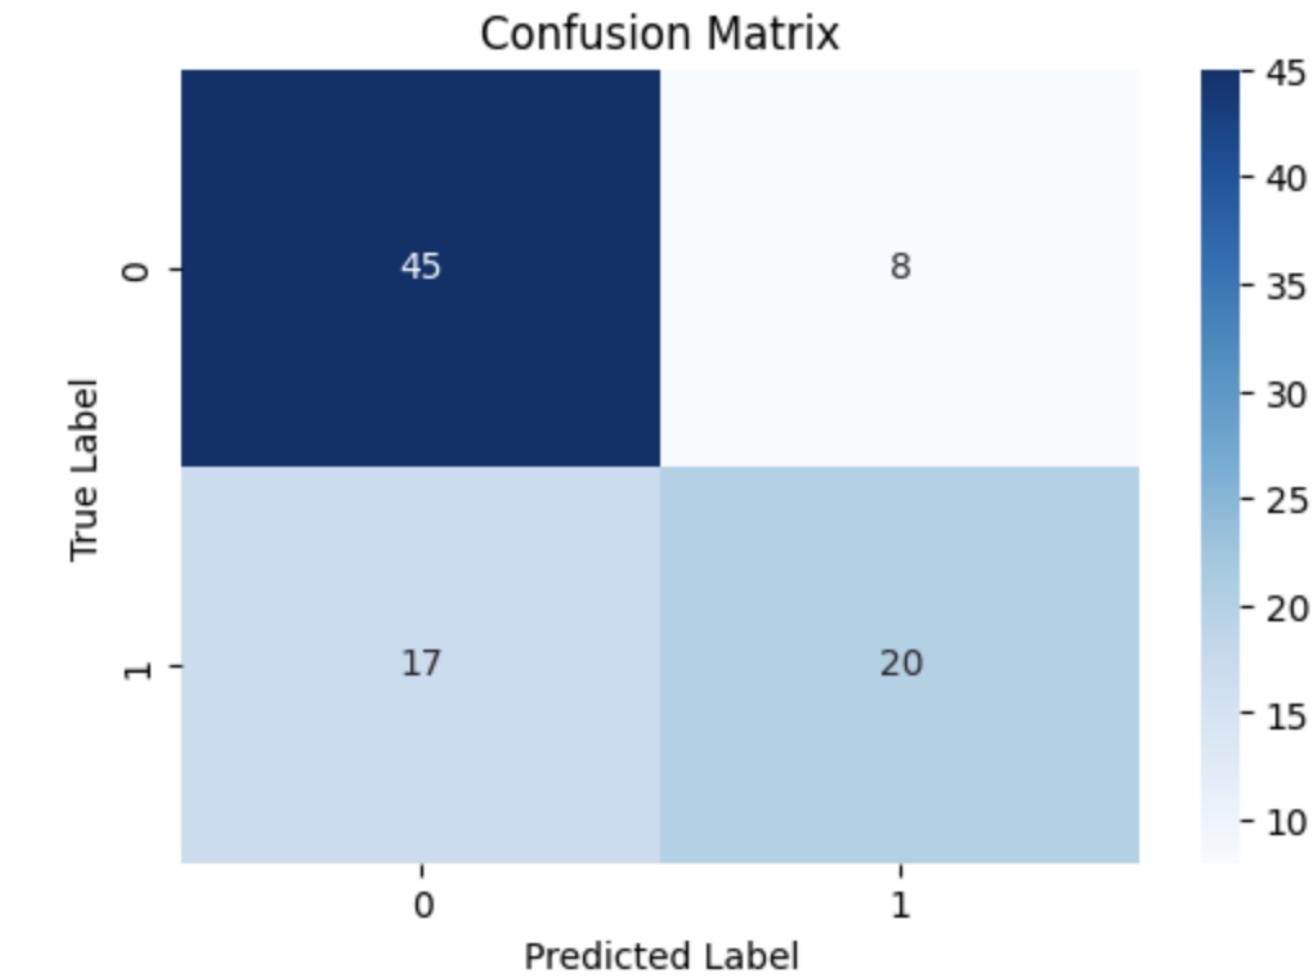
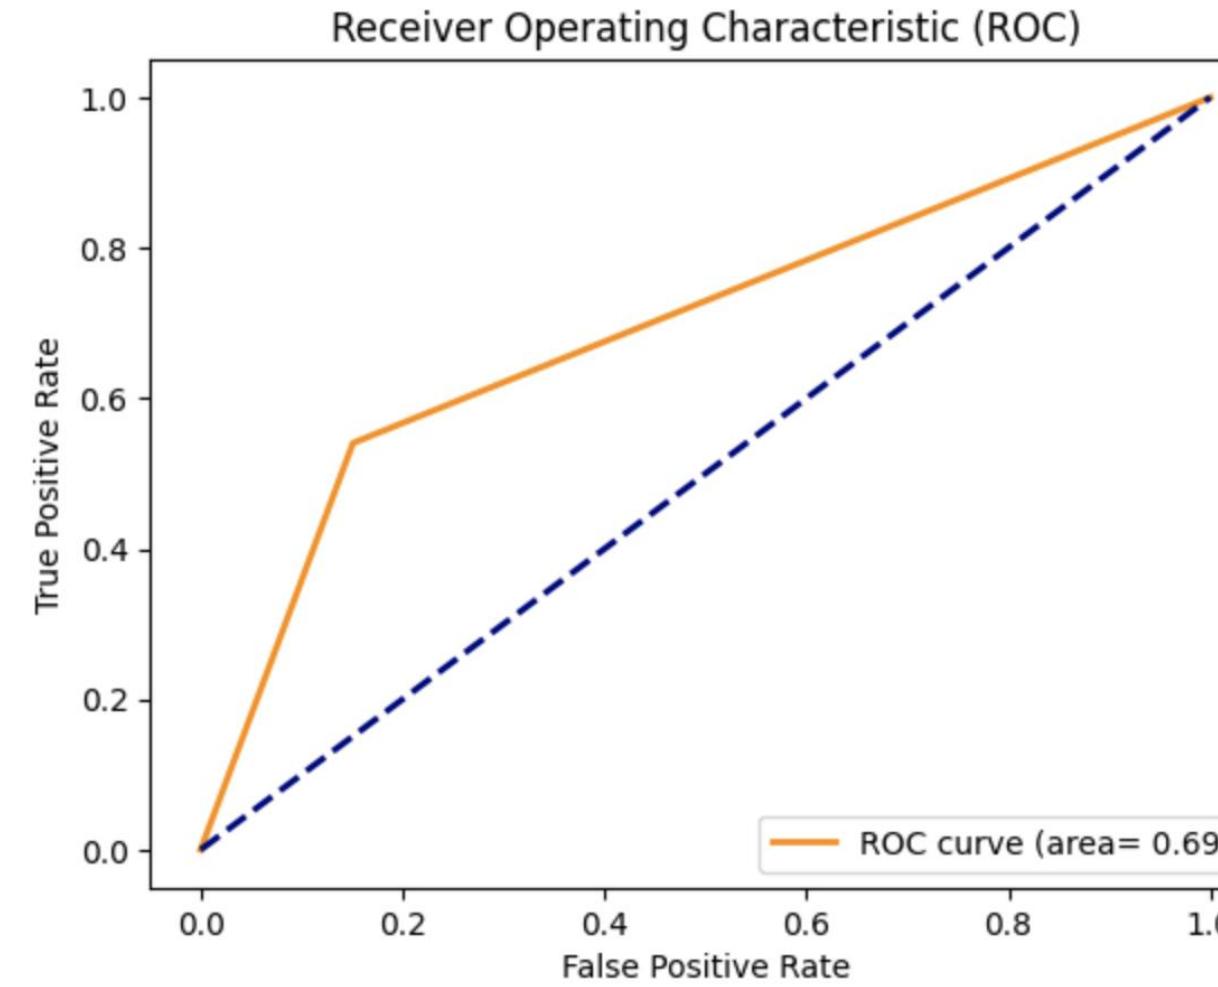
- Predicted outcomes (`y_pred`) on the **test set (`X_test`)**.

Model Evaluation:

- Calculated model accuracy: `{accuracy:.2f}` (replace with the printed accuracy).
- Generated a **classification report** to evaluate precision, recall, F1-score, and support for each class.

<Step 6>

<Model Evaluation>



- An ROC curve (area) of 0.69 indicates that the model is better than random guessing but still leaves room for improvement
- For the confusion matrix:
 - **False Negatives:** The model misses 17 "Death" cases, highlighting low recall (54.1%) for the positive class, which is critical in high-stakes contexts.
 - **False Positives:** There are 8 incorrect "Death" predictions; while less critical, they could lead to unnecessary actions.
 - **Class Imbalance:** Likely skewed towards "No Death," affecting the model's ability to detect minority cases.
 - **Model Simplicity:** A max depth of 3 limits the model's complexity, potentially causing misclassifications.

<Step 7>

<Model Interpretation and Insights>

ROC AUC Score (0.69):

- Moderate ability to discriminate between death and survival events.
- Not sufficient for reliable predictions in high-stakes healthcare settings.

Confusion Matrix Insights:

- 17 false negatives, indicating missed death events.
- Risks of not identifying patients needing critical intervention.

Key Features:

- **Serum creatinine** and **ejection fraction** align with medical knowledge of heart failure.
- Model captures some high-risk patterns but shows variability in **platelets** and **serum sodium** predictions.

Model Context:

- False negatives are critical in healthcare; missing death events can have severe consequences.
- Model conservatively identifies survival cases but risks overlooking critical patients.

Next Steps:

- Focus on reducing false negatives for better detection.
- Explore resampling or cost-sensitive methods to handle class imbalance.

<Thank You!>
Any Questions?