

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing

from google.colab import drive
drive.mount('/content/drive')

pip install statsmodels

↳ Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (0.14.4)
Requirement already satisfied: numpy<3,>=1.22.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (1.26.4)
Requirement already satisfied: scipy!=1.9.2,>=1.8 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (1.13.1)
Requirement already satisfied: pandas!=2.1.0,>=1.4 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (2.2.2)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (0.5.6)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (24.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2.36.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.2)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.6->statsmodels) (1.16.0)

```

```

from google.colab import drive
drive.mount('/content/gdrive')

```

**↳** Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```

#read data
path = "/content/gdrive/MyDrive/Colab Notebooks/My_Data/airline_passenger_satisfaction.csv"
df = pd.read_csv(path, index_col=0)
df.head()

```

ID	Gender	Age	Customer Type	Type of Travel	Class	Flight Distance	Departure Delay	Arrival Delay	Departure and Arrival Time Convenience	Ease of Online Booking	...	On-board Service	Seat Comfort	Leg Room Service	Cleanliness
1	Male	48	First-time	Business	Business	821	2	5.0	3	3	...	3	5	2	
2	Female	35	Returning	Business	Business	821	26	39.0	2	2	...	5	4	5	
3	Male	41	Returning	Business	Business	853	0	0.0	4	4	...	3	5	3	
4	Male	50	Returning	Business	Business	1905	0	0.0	2	2	...	5	5	5	
5	Female	49	Returning	Business	Business	3470	0	1.0	3	3	...	3	4	4	

5 rows × 23 columns

```
df.info()
```

**↳** <class 'pandas.core.frame.DataFrame'>  
Index: 129880 entries, 1 to 129880  
Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
0	Gender	129880	object
1	Age	129880	int64
2	Customer Type	129880	object
3	Type of Travel	129880	object
4	Class	129880	object
5	Flight Distance	129880	int64
6	Departure Delay	129880	int64
7	Arrival Delay	129487	float64
8	Departure and Arrival Time Convenience	129880	int64
9	Ease of Online Booking	129880	int64
10	Check-in Service	129880	int64
11	Online Boarding	129880	int64
12	Gate Location	129880	int64
13	On-board Service	129880	int64
14	Seat Comfort	129880	int64
15	Leg Room Service	129880	int64

```

16 Cleanliness           129880 non-null  int64
17 Food and Drink       129880 non-null  int64
18 In-flight Service    129880 non-null  int64
19 In-flight Wifi Service 129880 non-null  int64
20 In-flight Entertainment 129880 non-null  int64
21 Baggage Handling     129880 non-null  int64
22 Satisfaction          129880 non-null  object
dtypes: float64(1), int64(17), object(5)
memory usage: 23.8+ MB

```

## ✓ I/ Explore data analysis

```
#data dimension
df.shape
```

→ (129880, 23)

```
#check data types
#df.info()
```

```
#check missing value
df.isnull().sum()
```

	0
Gender	0
Age	0
Customer Type	0
Type of Travel	0
Class	0
Flight Distance	0
Departure Delay	0
Arrival Delay	393
Departure and Arrival Time Convenience	0
Ease of Online Booking	0
Check-in Service	0
Online Boarding	0
Gate Location	0
On-board Service	0
Seat Comfort	0
Leg Room Service	0
Cleanliness	0
Food and Drink	0
In-flight Service	0
In-flight Wifi Service	0
In-flight Entertainment	0
Baggage Handling	0
Satisfaction	0

```
# only Arrival Delay has 393 missing values, which is only a minor part of the data, so they are dropped
df.dropna(inplace=True)
```


```
#checking for imbalance of the target variable : Satisfaction
#convert to binary class for Satisfaction: neutral or dissatisfied(0), satisfied (1)
le = preprocessing.LabelEncoder()
```

```

le.fit(df.Satisfaction)
print(le.classes_)
df.Satisfaction = le.transform(df.Satisfaction)
#plot data
satis_count = df['Satisfaction'].value_counts('normalized')
ax = satis_count.plot(kind='pie', explode=(0.018,0), labels=['neutral or dissatisfied','satisfied'], autopct='%1.2f%%')
# set the title
ax.set_title('neutral or dissatisfied(0), satisfied (1)')

→ ['Neutral or Dissatisfied' 'Satisfied']
Text(0.5, 1.0, 'neutral or dissatisfied(0), satisfied (1)')
    neutral or dissatisfied(0), satisfied (1)

```



We can see that 43.45% of the customer satisfied with the flight service while the remain 56.55% tent to feel neutral or dissatisfied with the service

```


#correlation plot

# Selecting only numerical columns from the DataFrame
numerical_data = df.select_dtypes(include=['number'])

# Calculating the correlation matrix
correlation_matrix = numerical_data.corr()

# Plotting the correlation heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0, fmt='.2f', linewidths=0.5)
plt.title("Correlation Matrix of Numerical Data")
plt.show()

```



```
#Since Departure Delay and Arrival delay has high correlation score (0.97), we have to check for collinearity
```

```
#Method 1: Check by Z-test
```

```
from statsmodels.stats.weightstats import ztest
z_statistic, p_value = ztest(df['Departure Delay'], df['Arrival Delay'])
```

```
print(f"Z-statistic: {z_statistic}, P-value: {p_value}")
```

→ Z-statistic: -2.982374414813736, P-value: 0.002860219674711084

**P-value=0.00286 < 0.05**, we reject the null hypothesis and conclude that there is a statistically significant difference between the means of 'Departure Delay' and 'Arrival Delay'


```
#drop arrival delay
df = df.drop('Arrival Delay', axis=1)
```

The VIF results for "Departure Delay" and "Arrival Delay" indicate significant multicollinearity between these two variables, with VIF values far exceeding the common thresholds of 5 or 10. This suggests that these variables share a high degree of linear dependency, which could complicate the interpretation of their individual effects

```

#Explore object data
#Including 4 features: Gender, Customer Type, Class, Type of Travel
fig, ([ax1, ax2],[ax3, ax4]) = plt.subplots(2, 2, figsize=(18, 9))
# plot each pie chart in a separate subplot
ax1.pie(df['Gender'].value_counts('normalized'), autopct='%.1f%%', labels=df['Gender'].value_counts('normalized').index,startangle = 90)
ax1.set_title('Gender')
ax2.pie(df['Customer Type'].value_counts('normalized'), autopct='%.1f%%',labels=df['Customer Type'].value_counts('normalized').index,startangle = 90)
ax2.set_title('Customer Type')
ax3.pie(df['Class'].value_counts('normalized'), autopct='%.1f%%', labels=df['Class'].value_counts('normalized').index,startangle = 90)
ax3.set_title('Class')
ax4.pie(df['Type of Travel'].value_counts('normalized'), autopct='%.1f%%',labels=df['Type of Travel'].value_counts('normalized').index,startangle = 90)
ax4.set_title('Type of Travel')
plt.tight_layout()

```



```

#encode the data to convert text to integer
# Encode gender
le.fit(df['Gender'])
df['Gender'] = le.transform(df['Gender'])
# Encode customer type
le.fit(df['Customer Type'])
df['Customer Type'] = le.transform(df['Customer Type'])
# Encode class
le.fit(df['Class'])
df['Class'] = le.transform(df['Class'])
# Encode type of travel
le.fit(df['Type of Travel'])
df['Type of Travel'] = le.transform(df['Type of Travel'])

```

```

# explore numerical data
# create a figure with 4 subplots
fig, ((ax1, ax2, ax3, ax4)) = plt.subplots(1, 4, figsize=(12, 6))

sns.barplot(x=df['Departure and Arrival Time Convenience'].value_counts().index,
            y=df['Departure and Arrival Time Convenience'].value_counts().values, ax=ax1)
ax1.set_title('Departure and Arrival Time Convenience')

sns.barplot(x=df['Gate Location'].value_counts().index,
            y=df['Gate Location'].value_counts().values, ax=ax2)
ax2.set_title('Gate Location')

sns.distplot(df['Departure Delay'], ax=ax3)
ax3.set_title('Departure Delay')

sns.distplot(df['Flight Distance'], ax=ax4)
ax4.set_title('Flight Distance')

plt.tight_layout()
plt.show()

```

↳ <ipython-input-44-db39edd126e5>:14: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(df['Departure Delay'], ax=ax3)
<ipython-input-44-db39edd126e5>:17: UserWarning:

```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```


sns.distplot(df['Flight Distance'], ax=ax4)

```


Departure and Arrival Time Convenience




Gate Location



Departure Delay



Flight Distance



```
plt.figure(figsize=(10,10))
plt.subplot(2,1,1)
sns.kdeplot(df.loc[df["Satisfaction"]==1]["Age"],alpha=0.5,label="satisfied",shade=True)
sns.kdeplot(df.loc[df["Satisfaction"]==0]["Age"],alpha=0.5,label="neutral or dissatisfied",shade=True)
plt.title("KDE Plot of Age")
plt.legend()

plt.subplot(2,1,2)
sns.kdeplot(df.loc[df["Satisfaction"]==1]["Flight Distance"],alpha=0.5,label="satisfied",shade=True)
sns.kdeplot(df.loc[df["Satisfaction"]==0]["Flight Distance"],alpha=0.5,label="neutral or dissatisfied",shade=True)
plt.title("KDE Plot of Flight Distance")
plt.legend()
```

```

↳ <ipython-input-45-6cd257d74a55>:3: FutureWarning:
  `shade` is now deprecated in favor of `fill`; setting `fill=True`.
  This will become an error in seaborn v0.14.0; please update your code.


  sns.kdeplot(df.loc[df["Satisfaction"]==1]["Age"],alpha=0.5,label="satisfied",shade=True)
<ipython-input-45-6cd257d74a55>:4: FutureWarning:
  `shade` is now deprecated in favor of `fill`; setting `fill=True`.
  This will become an error in seaborn v0.14.0; please update your code.

  sns.kdeplot(df.loc[df["Satisfaction"]==0]["Age"],alpha=0.5,label="neutral or dissatisfied",shade=True)
<ipython-input-45-6cd257d74a55>:9: FutureWarning:
  `shade` is now deprecated in favor of `fill`; setting `fill=True`.
  This will become an error in seaborn v0.14.0; please update your code.

  sns.kdeplot(df.loc[df["Satisfaction"]==1]["Flight Distance"],alpha=0.5,label="satisfied",shade=True)
<ipython-input-45-6cd257d74a55>:10: FutureWarning:
  `shade` is now deprecated in favor of `fill`; setting `fill=True`.
  This will become an error in seaborn v0.14.0; please update your code.

  sns.kdeplot(df.loc[df["Satisfaction"]==0]["Flight Distance"],alpha=0.5,label="neutral or dissatisfied",shade=True)
<matplotlib.legend.Legend at 0x7f395a20ae00>

```



- Older customer seems easier on the review than the youth.
- Interesting point is that long flight customer satisfied more than short flight. Most of the bad review comes from flight distance less than 1500km. We could divide flight distance into short and and long distance flight instead of using the number.

```
#since both Departure Delay and Flight Distance are heavily right skewed, the data needs to be transformed
df['Departure Delay'] = np.log(df['Departure Delay'] + 1)
df['Flight Distance'] = np.log(df['Flight Distance'] + 1)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))

sns.distplot(df['Departure Delay'], ax=ax1)
ax1.set_title('Log-Transformed Departure Delay')
ax1.set_xlabel('Log Departure Delay')
ax1.set_ylabel('Density')

sns.distplot(df['Flight Distance'], ax=ax2)
ax2.set_title('Log-Transformed Flight Distance')
ax2.set_xlabel('Log Flight Distance')
ax2.set_ylabel('Density')

plt.tight_layout()
plt.show()
```

→ <ipython-input-46-ee7dd9501f2d>:8: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Departure Delay'], ax=ax1)
<ipython-input-46-ee7dd9501f2d>:13: UserWarning:
```


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Flight Distance'], ax=ax2)
```

Log-Transformed Departure Delay



Log-Transformed Flight Distance



```
from scipy.stats import chi2_contingency
```

```

# List of categorical columns, excluding 'Flight Distance', 'Departure Delay', 'Arrival Delay', and 'Satisfaction'
categorical_columns = [col for col in df.columns if col not in ['Flight Distance', 'Departure Delay', 'Satisfaction']]

# Dictionary to hold the p-values of the Chi-square tests
chi_square_results = {}

# Perform Chi-square test for each categorical variable against 'Satisfaction'
for col in categorical_columns:
    contingency_table = pd.crosstab(df[col], df['Satisfaction'])
    chi2, p, dof, expected = chi2_contingency(contingency_table)
    chi_square_results[col] = p

# Printing the p-values
for col, p in chi_square_results.items():
    print(f"Variable: {col}, P-value: {p}")

→ Variable: Gender, P-value: 3.608854694610591e-05
Variable: Age, P-value: 0.0
Variable: Customer Type, P-value: 0.0
Variable: Type of Travel, P-value: 0.0
Variable: Class, P-value: 0.0
Variable: Departure and Arrival Time Convenience, P-value: 2.473802119547831e-127
Variable: Ease of Online Booking, P-value: 0.0
Variable: Check-in Service, P-value: 0.0
Variable: Online Boarding, P-value: 0.0
Variable: Gate Location, P-value: 0.0
Variable: On-board Service, P-value: 0.0
Variable: Seat Comfort, P-value: 0.0
Variable: Leg Room Service, P-value: 0.0
Variable: Cleanliness, P-value: 0.0
Variable: Food and Drink, P-value: 0.0
Variable: In-flight Service, P-value: 0.0
Variable: In-flight Wifi Service, P-value: 0.0
Variable: In-flight Entertainment, P-value: 0.0
Variable: Baggage Handling, P-value: 0.0

```

Based on the p-values from the Chi-square tests, it is evident that all examined categorical variables have a statistically significant association with the 'Satisfaction' variable. This indicates that factors ranging from 'Gender' to 'Baggage Handling' potentially play a role in influencing customer satisfaction levels. The strength of these associations is robust, given the extremely low p-values, suggesting that these variables are important to consider when modeling customer satisfaction in logistic regression analysis.

## ✓ II/ Modeling

### a/ Prepare data for modelling

```

df = df.reset_index(drop=True)

from sklearn.model_selection import train_test_split
X = df.drop("Satisfaction", axis=1)
y = df["Satisfaction"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
print(X_train.shape)
print(X_test.shape)

→ (90640, 21)
(38847, 21)

```

### b/ Logistic Regression

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
lr = LogisticRegression(solver='liblinear')
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
print(lr)

```

```

    ➜ Accuracy: 0.8739671017066954
      precision    recall  f1-score   support
      0       0.88      0.90      0.89     21868
      1       0.87      0.84      0.85     16979

      accuracy                           0.87      38847
      macro avg       0.87      0.87      0.87      38847
      weighted avg    0.87      0.87      0.87      38847


LogisticRegression(solver='liblinear')

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Calculate the ROC curve and AUC for the model
fpr, tpr, thresholds = roc_curve(y_test, lr.predict_proba(X_test)[:,1])
roc_auc = auc(fpr, tpr)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```



c/ Decision Tree

```

from sklearn.tree import DecisionTreeClassifier,plot_tree # Import Decision Tree Classifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

tree = DecisionTreeClassifier(max_depth=4)
tree.fit(X_train,y_train)
y_pred= tree.predict(X_test)
print("Accuracy:",accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

```

```

    ➜ Accuracy: 0.8878420470049168
      precision    recall  f1-score   support
      0       0.88      0.93      0.90     21868
      1       0.90      0.84      0.87     16979


```

accuracy			0.89	38847
macro avg	0.89	0.88	0.89	38847
weighted avg	0.89	0.89	0.89	38847

Confusion Matrix:

```
[[20281 1587]
 [ 2770 14209]]
```

```
plt.figure(figsize=(20, 10))
plot_tree(tree, filled=True, class_names=['Dissatisfied or Neutral','Satisfied'])
plt.show()
```



```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

```
# Compute ROC curve and ROC area
```